

Recurrent Neural Network Language Modeling

Marco Gualtieri - Maurizio Gullo

Università degli Studi di Firenze

2014-2015

Language Modeling

A **LM** is a probability distribution over strings that reflects how frequently a string occurs as a sentence.

Goal: learn the joint probability function of sequences of words in a language, in order to predict the next word in textual data given context.

Applications:

- *speech recognition and handwriting recognition*
- *machine translation*
- *spell correction and intelligent input method*
- *information retrieval*

Problem

Data sparsity: just a very small subset of possible sentences will be observed during training.

Probabilistic LM

Compute the probability of a sentence:

$$P(W) = P(w_1, w_2, \dots, w_n)$$

by predicting probability of each upcoming word:

$$P(w_n | w_1, w_2, \dots, w_{n-1})$$

Basic idea:

Chain rule

$$\begin{aligned} P(w_1, w_2, \dots, w_n) &= P(w_1)P(w_2|w_1) \dots P(w_n|w_1, \dots, w_{n-1}) \\ &= \prod_i P(w_i | w_1, \dots, w_{i-1}) \end{aligned}$$

Probabilistic LM - estimate probs

Maximum Likelihood Estimate (frequency counts)

$$P(w_i | w_1, w_2, \dots, w_{i-1}) = \frac{C(w_1 w_2 \dots w_{i-1} w_i)}{C(w_1 w_2 \dots w_{i-1})}$$

This is not suitable, because there are too many possible sentences.

To simplify, we can use **Markov** assumption:

N-th order Markov assumption

$$P(w_i | w_1, w_2, \dots, w_{i-1}) \approx P(w_i | w_{i-N} \dots w_{i-1})$$

Depending on N we have *unigram*, *bigram* ... *N-gram*.

Probabilistic LM - improve generalization

With this model, if a word or a N-gram in the test set is not in the training set, we will give probability equal to 0 to the whole test set.

Possible improvements:

- **Smoothing**: assign some of the total probability mass to unseen words or N-grams
- **Back-off** models: recursively consider the back-off conditional probability of the (N-1)-gram if the N-gram has not been observed enough during training

Evaluate a LM

The best language model is the one that best predicts an unseen test set, giving to it the highest probability.

Maximizing the probability \iff Minimizing perplexity

Perplexity

$$PPL = P(w_1, w_2, \dots, w_n)^{-\frac{1}{n}} = 2^{-\frac{1}{n} \sum_{i=1}^n \log_2 p(w_i)}$$

PPL is the inverse probability of the test set, normalized by the number of words.

***log** space is convenient to avoid underflow.

Neural Network LM

N-gram and back-off limitations:

- short context
- no syntactic/semantic similarity


Bengio's NN introduces:

- word feature vector $M : V \rightarrow \mathbb{R}^m$
- $f(w_{i-C}, \dots, w_{i-1}, w_i) = g(M(w_{i-C}), \dots, M(w_{i-1}), M(w_i))$


Both mapping M and function g are learned.

In this way, similar words are expected to have similar feature vectors. Therefore during training each sentence will increase probability also for similar sentences.

Bengio's feed-forward NN



Recurrent Neural Network LM



$$x(t) = w_t + s(t - 1)$$

$$s_j(t) = f \left(\sum_i x_i(t) u_{ji} \right)$$

$$y_k(t) = g \left(\sum_j s_j(t) v_{kj} \right)$$

Activation functions

$$f(z) = \frac{1}{1 + e^{-z}}$$

$$g(z_m) = \frac{e^{z_m}}{\sum_k e^{z_k}}$$

RNN LM - Features

- history length can be unlimited
- the whole history is compressed in a low dimensional space (not only words)
- trained with a standard backpropagation with gradient descent
- only hidden layer size has to be set

It has been observed that RNN can form short-term memory in context layer neurons; while long-term memory resides in synapses.

So the network should continue training even during testing phase, using a fixed learning rate (**dynamic model**).

RNN - Softmax and error gradient

Softmax function forces values in output nodes to lie in $[0, 1]$ and to sum to 1, allowing to model a joint distribution over the output variables.

The derivative is:

$$\frac{\partial g_j}{\partial z_k} = g_j \delta_{kj} - g_j g_k$$

where δ_{kj} is the Kronecker Delta.

RNN - Softmax and error gradient

The **cross-entropy** function

$$H(p, q) = - \sum_x p(x) \log q(x)$$

measures similarity between distributions p (reference) and q (model).

Each data sample is drawn from V different classes, so our cross-entropy cost function is:

$$E = - \sum_{k=1}^V t_k \log g_k(z)$$

RNN - Softmax and error gradient

Given that the target t is constant, then minimizing the cross-entropy cost is equivalent to minimize this functional:

$$\tilde{E} = - \sum_{k=1}^V t_k \log g_k(z) + \sum_{k=1}^V t_k \log t_k = - \sum_{k=1}^V t_k \log \frac{g_k(z)}{t_k}$$

Considering that only one of the component of t is different from zero, the Jacobian of \tilde{E} takes this convenient form:

$$\frac{\partial \tilde{E}}{\partial z_j} = g_j - t_j \quad \forall j = 1, \dots, V$$

Implementation notes - Computation reductions

In the feedforward phase, we have to compute the dot product between:

- input vector x , with length $(V + H)$
- matrix W_1 , with size $(V + H) \times H$

Result is a vector with length H .

Since the first V values of x are in 1-of- N coding, we can compute the dot product between the last H elements of x and the last H rows of W_1 , then sum the k^{th} row of W_1 to the result (k is the index of the current word).

The same approach can be applied when backpropagating the error from the hidden layer to the input.

Implementation notes - Stopping criterion

At the end of each epoch, we evaluate PPL on a validation set:

- If this has increased compared to the previous iteration, we restore the weights of the network to the values at the previous step.
- Whether this value is increased or just reported a not significant improvement, we halve the value of the learning rate and keep on halving it in the following epochs.
- When the "halving condition" is encountered again, we stop the execution of the algorithm.

Convergence is usually achieved in 10/20 epochs.

Implementation notes - Others

- scripts to generate vocabulary (with frequency of each word) and dataset, cleaning punctuation and special chars
- recovery system at each epoch (to restart learning from the last one)
- momentum

Packages:

- *Python* and *NumPy*
- *openblas* on *Ubuntu*
- *Anaconda Accelerate* and *MKL* on *Mac OSX*

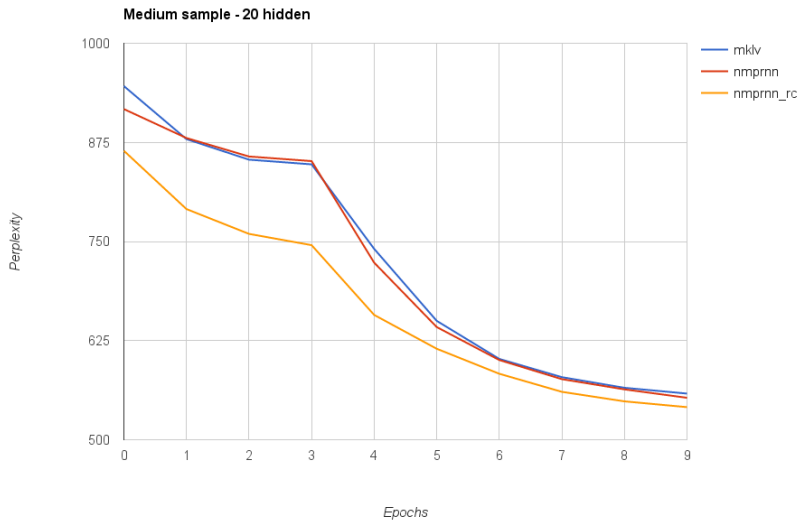
Some results

Comparison on medium dataset: 600K words, 25K tokens


	Epochs	Validation PPL	Test PPL
mklv_20	10	558.2	602.9
nmprnn_20	10	552.9	603.7
nmprnn_20_rc	10	541.1	594.4
mklv_100	9	579.9	642.5
nmprnn_100	10	546.4	607.6
nmprnn_100_rc	9	546.5	602.0
mklv_200	9	629.5	690.3
nmprnn_200	10	590.2	662.6
nmprnn_200_rc	9	569.2	640.4

***rc**: resetting of the context layer at each new sentence


Some results



Some results



Some results



Complexity of training

I : number of training epochs for convergence

W : number of tokens in training set

N : N-gram order

D : size of each word's feature vector

H : size of hidden layer

V : size of vocabulary

N-gram NN:

$$I \times W \times ((N - 1) \times D \times H + H \times V)$$

RNN:

$$I \times W \times (H \times H + H \times V)$$

Improvements - Operations on data

Reduction of tokens

Idea: train the network with *in-domain* data (usually 10% of the corpus) and a subsample of the *out-of-domain* data.

Bad subsampling can lead to results degradation.

Automatic selection and sorting

Idea: split training set into chunks and sort them by perplexity (computed with bigrams), discarding too noisy data.

Incremental learning: let the network processing *out-of-domain* data first (basic pattern) and most important *in-domain* data at the end (complex pattern). Data processed at the end will have higher weights.

Improvements - Reduction of $H \times V$ term

Reduce hidden layer size

To obtain good results, this can't be reduced too much.

Reduce vocabulary size

Goodman's idea: assign each word to a class, and compute probability distribution over classes only (with RNN).

Then compute distribution over words by assuming unigram distribution of words within a class. With C classes:

$$I \times W \times (H \times H + H \times C)$$

Compression layer


Idea: insert a compression layer (size P) between hidden and output:

$$I \times W \times (H \times H + H \times P + P \times V)$$

Improvements - BPTT

Backpropagation **through time** is a gradient-based technique to train RNN.

The idea is to *unfold* the network through time when feedforwarding, like shown in figure (here time depth is 3):



then backpropagate the error as usual, averaging the weights of each instance of f to compute \mathbf{x}_{t+1} .

BPTT tends to be faster for training RNN, but has problems with local minima.

Improvements - Interpolation

Interpolate different models (even of a different type, like a RNN and a N-gram) usually leads to better results:

$$P(w_n | w_1, \dots, w_{n-1}) \approx \lambda_1 \text{model}_1(w_1, \dots, w_n) + \dots + \lambda_k \text{model}_k(w_1, \dots, w_n)$$

where


$$\sum_{i=1}^k \lambda_i = 1$$

Usually lambdas are chosen to maximize the probability of a validation/test set.

Deficiencies of backpropagation with gradient-descent

- poor generalization with many hidden layers (deep networks)
- slow progress on deep nets
- early stopping before making significant progress (under-fitting)
- is unsuitable for objectives that exhibit "pathological curvature"

Pathological curvature



Newton's method

General idea is that f can be approximated around each θ by the quadratic:

$$f(\theta + p) \approx q_\theta(p) \equiv f(\theta) + \nabla f(\theta)^\top p + \frac{1}{2} p^\top B p$$

where $B = H(\theta)$ is the Hessian matrix of f at θ .

Goal

Find a good search direction minimizing this quadratic with respect to p

In the standard Newton's method, the optimal p is obtained by computing the $N \times N$ matrix B and then solving the system $Bp = -\nabla f(\theta)$.

This is prohibitively expensive when N is large, as it is in usual neural networks.

Hessian-free optimization

Hessian-free optimization uses two simple ideas:

- 1 Given an N -dimensional vector d , Hd can be easily computed using finite differences at the cost of a single extra gradient evaluation:

$$Hd = \lim_{\epsilon \rightarrow 0} \frac{\nabla f(\theta + \epsilon d) - \nabla f(\theta)}{\epsilon}$$

- 2 An effective algorithm like *conjugate gradients* can be used for optimizing quadratic objectives (such as $q_\theta(p)$).

Hessian-free optimization method

Algorithm 1 Hessian-free optimization method

```
1: for  $n = 1, 2, \dots$  do  
2:    $g_n \leftarrow \nabla f(\theta_n)$   
3:   compute/adjust  $\lambda$  by some method  
4:   define function  $B_n(d) = H(\theta_n)d + \lambda d$   
5:    $p_n \leftarrow \text{CG-minimize}(B_n, -g_n)$   
6:    $\theta_{n+1} \leftarrow \theta_n + p_n$   
7: end for
```

References

A Neural Probabilistic Language Model

Bengio Y., Ducharme R., Vincent P., Jauvin C.

Université de Montréal, Département d'Informatique et Recherche Opérationnelle
Journal of Machine Learning Research, 2003.

Recurrent neural network based language model

Mikolov T., Karafiat M., Burget L., Cernocky J., Khudanpur S.

Brno University of Technology

Interspeech, 2010.

Strategies for Training Large Scale Neural Network Language Models

Mikolov T., Deoras A., Povey D., Burget L., Cernocky J.

Brno University of Technology

Proceedings of ASRU, 2011.

Extensions of Recurrent Neural Network Language Model

Mikolov T., Kombrink S., Burget L., Cernocky J., Khudanpur S.

Brno University of Technology

ICASSP, 2011.

References

Recurrent neural network language modeling toolkit

Mikolov T., Kombrink S., Deoras A., Burget L., Cernocky J.

Brno University of Technology

ASRU Demo Session, 2011.

Deep learning via Hessian-free optimization

Martens J.

University of Toronto

ICML, 2010.

Classes for fast maximum entropy training

J. Goodman

Microsoft Research, Redmond

ICASSP, 2001.

github.com/marcogualtieri/rnnlm_numpy